

## **CPU Project Description**

This project involves building a high-level simulator for a simple 64-bit RISC CPU in C. The instruction set architecture (ISA, defining the registers, instructions, etc.) of this CPU is based on an early version of the MIPS processor (which was influential in its day), except that only integer registers and instructions will be simulated and various other details (floating point registers and instructions, overflow, etc.) will be left out. Additionally, the MIPS processor was a 32-bit processor, so I have taken the liberty of modifying the registers and instruction set to be 64 bits.

### **The MIPS Registers**

There are 32 integer registers, numbered from 0 to 31. Each register is 64-bits. Register 0 always contains the value 0 and cannot be modified. Other of these registers are used for special purposes (stack pointer, base pointer, etc.), but that is by convention and does not affect the implementation of your CPU. In this document, we will refer to the  $i^{\text{th}}$  register as  $R[i]$ .

Finally, there is a program counter register, PC, which is not explicitly referenced in any instruction (we have referred to the program counter in class as an “instruction pointer” register).

### **The MIPS Instructions**

All instructions are exactly 64 bits wide.

There are three categories of instructions, R-instructions, I-instructions, and J-instructions. Within each category, every instruction has the same format, as described below.

#### **R-Instructions**

The R-instructions are used to perform operations where the operands are registers.

R-instructions have the following format, from left to right (most significant bit to least significant bit):

- opcode field: 6 bits (NOTE: for all R-instructions, opcode = 0)
- RS field: 5 bits (specifies one of the registers)
- RT field: 5 bits (specifies one of the registers)
- RD field: 5 bits (specifies one of the registers)
- UNUSED: 31 bits (ignored)
- shamt field: 6 bits (specifies a shift amount in a shift instruction)
- funct field: 6 bits (specifies which R-instruction to execute)

The list of R-instructions that you will need to implement, and their descriptions, is as follows.

(Remember, the opcode for R-instructions is always 0)

Instr	Description	Funct	Operation	Note
add	signed addition	0x20	$R[rd] \leftarrow R[rs] + R[rt]$	
sub	signed subtraction	0x22	$R[rd] \leftarrow R[rs] - R[rt]$	
mult	signed multiplication	0x18	$R[rd] \leftarrow R[rs] * R[rt]$	
div	signed division	0x1a	$R[rd] \leftarrow R[rs] / R[rt]$	
mod	modulo	0x1b	$R[rd] \leftarrow R[rs] \% R[rt]$	“%” is the mod operator in C
and	bitwise AND	0x24	$R[rd] \leftarrow R[rs] \& R[rt]$	
or	bitwise OR	0x25	$R[rd] \leftarrow R[rs]   R[rt]$	
xor	bitwise XOR	0x26	$R[rd] \leftarrow R[rs] \wedge R[rt]$	
nor	bitwise NOR	0x27	$R[rd] \leftarrow \sim(R[rs]   R[rt])$	
slt	set less than (signed)	0x2a	If $(R[rs] < R[rt])$ $R[rd] \leftarrow 1$ else $R[rd] \leftarrow 0$	The comparison treats $R[rs]$ and $R[rt]$ as signed numbers.
sltu	set less than (unsigned)	0x2b	If $(R[rs] < R[rt])$ $R[rd] \leftarrow 1$ else $R[rd] \leftarrow 0$	The comparison treats $R[rs]$ and $R[rt]$ as unsigned numbers.
sll	shift left logical	0x00	$R[rd] \leftarrow R[rt] \ll \text{shamt}$	
sllv	shift left logical variable	0x04	$R[rd] \leftarrow R[rs] \ll R[rt]$	
srl	shift right logical	0x02	$R[rd] \leftarrow R[rt] \gg \text{shamt}$	
srlv	shift right logical variable	0x06	$R[rd] \leftarrow R[rs] \gg R[rt]$	
sra	shift right arithmetic	0x03	$R[rd] \leftarrow R[rt] \gg \text{shamt}$	See discussion below about arithmetic shift.
srav	shift right arithmetic variable	0x07	$R[rd] \leftarrow R[rs] \gg R[rt]$	See discussion below about arithmetic shift.
jr	jump register	0x08	$\text{npc} \leftarrow R[rs]$	The target of the jump is specified in $R[rs]$ . See discussion below concerning npc and pc.
jalr	jump and link register	0x09	$r[31] \leftarrow \text{pc} + 8$ $\text{npc} \leftarrow r[rs]$	For procedure calls. The return address ( $\text{pc}+8$ ) is stored in $R[31]$ . The target of the jump is specified in $R[rs]$ . See discussion below concerning npc and pc.
syscall	System call for OS services	0x0c		See discussion below concerning syscall.

### **I-instructions**

I-instructions are used to perform operations using an immediate value (i.e. a constant). These include some arithmetic instructions, load and store instructions, and conditional branch instructions.

I-instructions have the following format, from left to right (most significant bit to least significant bit):

- opcode field: 6 bits (specifies which instruction to execute)
- RS field: 5 bits (specifies one of the registers)
- RT field: 5 bits (specifies one of the registers)
- immediate field: 48 bits (specifies a constant value to be used)

The list of I-instructions that you will need to implement, and their descriptions, is as follows.

<b>Instr</b>	<b>Description</b>	<b>opcode</b>	<b>Operation</b>	<b>Note</b>
addi	add immediate	0x08	$R[rt] \leftarrow R[rs] + \text{sign\_extend}(\text{imm})$	See discussion on sign extension below
addiu	add immediate unsigned	0x09	$R[rt] \leftarrow R[rs] + \text{imm}$	
lw	load word	0x23	$R[rt] \leftarrow M[R[rs] + \text{sign\_extend}(\text{imm})]$	$M[]$ represents memory.
lh	load half-word	0x21	$R[rt] \leftarrow \text{sign\_extend}(M[R[rs] + \text{sign\_extend}(\text{Imm})])$	32 bits are loaded from memory into the lower half of $R[rt]$ . The upper half of $R[rt]$ results from sign-extending the loaded value.
lhu	load half-word unsigned	0x25	$R[rt] \leftarrow M[R[rs] + \text{sign\_extend}(\text{Imm})]$	32 bits are loaded from memory into the lower half of $R[rt]$ . The upper half of $R[rt]$ contains only zeros (no sign extension). Note that the immediate operand is sign extended from 48 to 64 bits, so can be a negative number.
lb	load byte	0x20	$R[rt] \leftarrow \text{sign\_extend}(M[R[rs] + \text{sign\_extend}(\text{Imm})])$	One byte is loaded from memory into the lowest byte of $R[rt]$ . The rest of $R[rt]$ results from sign-extending the loaded value.

Instr	Description	opcode	Operation	Note
lbu	load byte unsigned	0x24	$R[rt] \leftarrow M[R[rs] + \text{sign\_extend}(\text{Imm})]$	One byte is loaded from memory into the lowest byte of $R[rt]$ . The rest of $R[rt]$ contains only zeros. Note that the immediate operand is sign-extended.
sw	store word	0x2b	$M[R[rs] + \text{sign\_extend}(\text{Imm})] \leftarrow R[rt]$	
sh	store half-word	0x29	$M[R[rs] + \text{sign\_extend}(\text{Imm})] \leftarrow R[rt]$	The lower half (32 bits) of $R[rt]$ is stored in memory.
sb	store byte	0x28	$M[R[rs] + \text{sign\_extend}(\text{Imm})] \leftarrow R[rt]$	The lowest byte of $R[rt]$ is stored in memory.
lui	load upper immediate	0x0f	$R[rt] \leftarrow (\text{Imm} \ll 32)$	Imm is loaded into the upper half of $R[rt]$ . The lower half of $R[t]$ contains only zeros. Note that since Imm is 48 bits, the upper 16 bits of Imm will be lost, which is only a problem if it contains a number too large to fit in 32 bits.
ori	bitwise or immediate	0x0d	$R[rt] \leftarrow R[rs] \mid \text{Imm}$	Imm is zero extended (i.e. the upper 16 bits are zero).
andi	bitwise and immediate	0x0c	$R[rt] \leftarrow R[rs] \& \text{Imm}$	Imm is zero extended.
xori	bitwise xor immediate	0x0e	$R[rt] \leftarrow R[rs] \wedge \text{Imm}$	Imm is zero extended.
slti	set less than immediate	0x0a	if $(R[rs] < \text{sign\_extend}(\text{imm}))$ $R[rt] \leftarrow 1$ else $R[rt] \leftarrow 0$	
sltiu	set less than immediate unsigned	0x0b	if $(R[rs] < \text{imm})$ $R[rt] \leftarrow 1$ else $R[rt] \leftarrow 0$	The comparison is unsigned, so both operands are interpreted as non-negative. Imm is not sign extended.

<b>Instr</b>	<b>Description</b>	<b>opcode</b>	<b>Operation</b>	<b>Note</b>
beq	branch on equal	0x04	if( $R[rs] == R[rt]$ ) $npc \leftarrow pc + 8 + \text{sign\_extend}(Imm)$	See discussion below on npc and pc.
bne	branch on not equal	0x05	if( $R[rs] != R[rt]$ ) $npc \leftarrow pc + 8 + \text{sign\_extend}(Immediate)$	

### **J-instructions**

J-instructions are jump instructions.

J-instructions have the following format.

- opcode field: 6 bits (specifies which instruction to execute)
- UNUSED: 10 bits (ignored)
- offset field: 48 bits (used in computing the target address of the jump)

The list of J-instructions that you will need to implement, and their descriptions, is as follows:

<b>Instr</b>	<b>Description</b>	<b>Opcode</b>	<b>Operation</b>	<b>Note</b>
j	jump	0x02	$npc \leftarrow (pc+8) + \text{sign\_extend}(offset)$	The target of jump comes from adding the sign-extended offset to pc+8 (i.e. the address of the following instruction)
jal	jump and link	0x03	$R[31] \leftarrow pc+8$ $npc \leftarrow (pc+8) + \text{sign\_extend}(offset)$	For procedure calls. The return address (pc+8) is stored in R[31]. The target of jump comes from the sign-extended offset to pc+8 (i.e. the address of the following instruction)

### **Implementing the MIPS Simulator**

I have provided you with a number of files, as follows:

- cpu.h: Declares the global variables representing the registers and memory, defines related constants, and declares the procedures needed to initialize the CPU, execute the CPU, and stop the CPU.
- cpu.o: The compiled version of cpu.c, which defines the variables and procedures declared in cpu.h as well as implementing the core CPU operation, as described below.
- r\_instructions.h: Declares the procedures that implement the R-instructions (other than syscall), one procedure per instruction.
- r\_instructions.o: The compiled version of r\_instructions.c, which defines the procedures declared in r\_instructions.h.

- `i_instructions.h`: Declares the procedures that implement the I-instructions, one procedure per instruction.
- `i_instructions.o`: The compiled version of `i_instructions.c`, which defines the procedures declared in `i_instructions.h`.
- `j_instructions.h`: Declares the procedures that implement the J-instructions, one procedure per instruction.
- `j_instructions.o`: The compiled version of `j_instructions.c`, which defines the procedures declared in `j_instructions.h`.
- `syscall.c`: This defines the procedure representing the syscall instruction.
- `Makefile`: Allows for the compilation of the `.c` files in the project, as necessary, by simply typing “make” in the shell. It generates an executable file name “cpu” (or `cpu.exe` on Windows), which you can run by typing “./cpu”.
- `test_programs.c`: Contains the `main()` procedure as well as procedures for testing the CPU on MIPS machine code programs. You should probably not modify this code unless you really understand what’s going on inside it. The `main()` procedure causes instructions to be loaded into memory starting at address `CODE_STARTING_ADDRESS` (see `cpu.h`).
- `utilities.h`, `utilities.c`: Contains procedures for printing out the bits of a 64-bit integer. You may find this useful for debugging.
- `register_names.h`, `programming.h`: Code I wrote to aid in testing the CPU (used in `test_programs.c`). You can take a look at these files, but you won’t need to use them.
- `ben_cpu`: This is a working version of the MIPS simulator that you can compare the output of your version to. To run it, simply type “./ben\_cpu”.

The programming project comprises three parts:

Part 1: Implement your own version of `r_instructions.c`.

Part 2: Implement your own versions of `i_instructions.c` and `j_instructions.c`.

Part 3: Implement your own version of `cpu.c`.

### **The use of “standardized” integer sizes**

As discussed in class, C has traditionally used the integer types `char`, `short`, `int` and `long`, as well as unsigned variants of them. As machines have evolved from 16-bit words to 32-bit words to the current 64-bit words, the sizes of `short`, `int`, and `long` have changed (`char` is always one byte, 8 bits). In order to ensure portability of C programs between different machines, programmers in industry tend to use standardized integer types with explicit sizes, defined in the `stdint.h` library. Therefore, we will also use these standardized integer types, rather than using `int`, `unsigned long`, etc. These standardized integer types are:

- `int64_t` (64-bit signed integer)
- `uint64_t` (64-bit unsigned integer)
- `int32_t` (32-bit signed integer)
- `uint32_t` (32-bit unsigned integer)
- `int16_t` (16-bit signed integer)
- `uint16_t` (16-bit unsigned integer)
- `int8_t` (8-bit signed integer)
- `uint8_t` (8-bit unsigned integer)

To use these standardized integer types, please note:

1. In any file that uses these types, put

```
#include <stdint.h>
```

at or near the top of the file (before the `#include` for any `.h` file from this project).

2. When using `printf` or `scanf` to print or read one of the 64-bit integer types, above, use “%lld” or “%llu” as the format specifier. The “ll” refers to “long long”, which is an old name for a 64-bit integer (when a `long` was 32 bits). These days, both `long long` and `long` are 64 bits, but the compiler still complains when “%ld” or “%lu” is used with an `int64_t` or `uint64_t`.

### **The simulated integer registers**

As you can see in `cpu.h`, the integer registers are declared as follows:

```
uint64_t registers[];
```

The procedures that you write to implement the various MIPS instructions should use these registers as appropriate. Thus, even though this document refers to the  $i^{\text{th}}$  register as `R[i]`, your code would refer to it as `registers[i]`.

### **The simulated memory**

Memory is declared as an array of bytes (not words) as follows:

```
uint8_t memory[];
```

Any instruction that you implement that accesses memory (e.g. lw, lh, lb, sw, sh, sb, etc.) will need to access this array. Be sure that your code is loading or storing the right number of bytes. For example, lw should load a 4-byte quantity from memory into a register. Referencing `memory[...]` in your code will just load or store a single byte, so to load a 4-byte (32-bit) quantity from memory, use a pointer of type `(uint32_t *)` to point to the area of `memory` that you want to load from or store to. For example, if you want to copy 4 bytes from memory into a variable `x` of type `uint32_t`, you could write something like:

```
x = *((uint32_t *) &memory[addr]);
```

### **The program counter (pc and npc)**

For convenience in the code, there are two variables that implement the program counter.

```
uint32_t pc;
uint32_t npc;
```

The variable `pc` always points to the current instruction being executed and `npc` is assigned to the *next* instruction to be executed. That is, the main loop of the CPU (as implemented by `cpu_execute()` in `cpu.c`) is:

*loop while not done:*

*pc ← npc*

*npc ← pc + 4*

*execute instruction at memory[pc]*

*end loop*

This way, if the executed instruction does not modify `npc`, the next instruction in memory will be executed in the next iteration of the loop. If the instruction being executed is a jump or branch instruction, it should set `npc` to the address to jump to. **IMPORTANT:** instructions should not modify `pc`, they should only read the value of `pc`. Conversely, `npc` may be written to (but only by jump or branch instructions) and should not be read. No other instructions should write to `npc`.

### **Sign Extension**

There are a number of instructions (e.g. most I-instructions) that require taking a 48-bit number and extending it to 64 bits. If the 48-bit number is non-negative, i.e. if bit 47 (the leftmost bit) is zero, then the leftmost 16 bits of the extended (64-bit) version of the number should be all zeros. But, if the 48-bit number is negative, i.e. if bit 47 is one, then the leftmost 16 bits of the extended version of the number should be all ones. For example,

#### Sign-extending a positive 48-bit number to a 64-bit number

```
0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 1101  ⇒
0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 1101
(this is 13 in decimal)
```

#### Sign-extending a negative 48-bit number to a 64-bit number

```
1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1101  ⇒
1111 1111 1111 1111 1111 1111 1111 1101 1111 1111 1111 1111
(this is -3 in decimal)
```



To implement this sign extension, perform the usual masking to check the value of the leftmost bit of the 47-bit value and, if that bit is not 0, then use a mask and the bitwise-OR operation to write all 1's in the upper 16 bits of the result.

The same holds for sign-extending an 8-bit number or a 32-bit number to a 64 bit number.

**NOTE:** You could implement the same sign-extension by relying on the C compiler, but I'd like you to use masking as described above (so you will know about sign extension if I ask it on the final exam).

### Arithmetic Right Shift

In order to use right shift as division by 2 or greater powers of two, it is necessary to preserve the sign of the number being shifted. To do so, the leftmost bit of the number before shifting should be repeated in the bits that opened up as a result of the shift. For example,

#### Shifting right by 3

<u>Before shift</u>	$\Rightarrow$	<u>After shift</u>
0011 0001 0010 ...		0000 0110 0010 010 ...
1011 0001 0010 ...		1111 0110 0010 010 ...

Here are some hints for implementing an arithmetic shift to the right by  $n$  bits:

- Perform the usual masking to test if the leftmost bit (bit 63) of the original number is 0 or not.
- If the leftmost bit of the original number is not 0, then after you shift the number to the right by  $n$ , you'll need to use a mask write  $n$  ones into the leftmost  $n$  bits of the shifted number.
- To create a mask with  $n$  ones in the rightmost position, use  $(1 \ll n) - 1$ . This, of course, is  $2^n - 1$ .
- That mask will need to be shifted to the left by  $(64 - n)$  to put the ones in the leftmost position.

**NOTE:** You could implement the same sign-extension by relying on the C compiler, but I'd like you to use masking as described above (so you will know about arithmetic shift right if I ask it on the final exam).

### Syscall

The syscall ("system call") instruction is an R-instruction, with no operands, that traps to the operating system so that the OS can perform a requested operation. You'll learn much more about system calls in your Operating Systems course. If you look in the `r_instructions.h` file, you'll see that the syscall procedure is declared by

```
void syscall(uint64_t instruction);
```

**I have already defined the syscall() procedure for you in syscall.c**, but looking at this code may be helpful to you. All your `r_instruction` code has to do is call `syscall()` when the syscall instruction is encountered.

The code I wrote for `syscall()` in `syscall.c` operates as follows. It exams the value stored in `R[2]` (which is the register whose nickname is `$V0`). Depending on the value of `R[2]`, `syscall()` perform the following actions:

`R[2] == 1`: (“`print_int`”) prints the integer value stored in `R[4]` (also known as `$A0`)  
`R[2] == 2`: (“`print_float`”) not supported, prints an error message and call `cpu_exit(1)`  
`R[2] == 3`: (“`print_double`”) not supported, prints an error message and call `cpu_exit(1)`  
`R[2] == 4`: (“`print_string`”) prints the string that starts at `Mem[R[4]]`.  
`R[2] == 5`: (“`read_int`”) read in an integer and store it in `R[2]` (also known as `$V0`)  
`R[2] == 6`: (“`read_float`”) not supported, prints an error message and call `cpu_exit(1)`  
`R[2] == 7`: (“`read_double`”) not supported, prints an error message and call `cpu_exit(1)`  
`R[2] == 8`: (“`read_string`”) read a string into memory starting at `Mem[R[4]]`. `R[5]` contains the size of the string buffer, so make sure the string is no longer than `R[5]` bytes.  
`R[2] == 9`: (“`sbrk`”) not supported, prints an error message and call `cpu_exit(1)`  
`R[2] == 10`: (“`exit`”) call `cpu_exit(0)`  
Otherwise: unsupported operation, prints an error message and call `cpu_exit(1)`

**Again, you do not have to implement anything in `syscall.c`.**

Post on the discussion board if you have any questions. Good luck!

**\*\*\*\* IMPORTANT: See the next page for hints! \*\*\*\***

## Hints

1. This assignment is not too hard, so don't stress out. It is long, though, so get started right away.
2. Use **this** document, **not** online materials, for understanding how the MIPS instructions should work.
3. You'll be working primarily with unsigned 64-bit integers, `uint64_t`, and signed 64-bit integers, `int64_t`. You will need to be careful when writing shifts of constant values, such as

```
(1 << 47)
```

because the compiler assumes that an integer literal, such as the 1 above, is a 32-bit integer. Thus, it will complain that the shift amount, 47, is too large. To ensure that you are generating a 64-bit value, you should cast the integer literal to a 64-bit number by using a cast:

```
((uint64_t) 1 << 47)
```

For example, to define a mask to check the value of bit 63 of a 64-bit number, you might write:

```
#define BIT63_MASK ((uint64_t) 1 << 63)
```

With any macro defined using `#define`, be sure to wrap the value in parentheses, as discussed in class.

4. The first file you need to implement is `r_instructions.c`. Logically, there are two parts to implementing the R-instructions:
  - Write macros (using `#define`) to extract the various fields of an R-instruction, i.e. the opcode, RS, RT, RD, shamt, and funct fields. These macros should use shifts and masks, just like the macros you wrote to extract the sign, exponent, and fraction fields of floating point numbers in assignment 2.
  - Implement each R-instruction (add, sub, etc.) as a separate, tiny function. The complete list of the functions you need to implement is found in `r_instructions.h`. To get you started, the `add()` function, which takes the 32-bit instruction as a parameter, should perform the following operations:
    - Extract the RS, RT, and RD fields from the instruction (the other fields aren't needed).
    - Compute the sum of `registers[RS]` and `registers[RT]` and write the result into `registers[RD]`.

Since `registers` is an array of `uint64_t`, but the add instruction is signed addition, you'll need to cast `registers[RS]` and `registers[RT]` to `int64_t`, perform the addition, and then cast the result back to `uint64_t` before writing it to `registers[RD]`. This function only needs to be a few (between one and four) lines.

5. The process for implementing `i_instructions.c` and `j_instructions.c` is much the same as for `r_instructions.c`, namely defining the macros (e.g. using masks) you need and then defining the very simple functions listed in `i_instructions.h` and `j_instructions.h`. Be sure to pay close attention to how each I-instruction and J-instruction is described in this document. Remember that each J-instruction (which is a jump) should write the computed destination address to `npc`.
6. In the last part of the assignment, Part 3, which is implementing `cpu.c`, you will be providing the code for several functions, including `cpu_execute()`. To implement `cpu_execute()`, you should think about how to call the appropriate function in `r_instructions.c`, `i_instructions.c`, or `j_instructions.c`, based on the opcode of an instruction. For example, if the opcode is `0x08`, the `addi()` function (see `r_instructions.h`) should be called.

Furthermore, if the opcode is 0, a function in `r_instructions.c` should be called based on the `funct` field. For example if the opcode is 0 and the `funct` field is `0x18`, then the `mult()` function (see `r_instructions.h`) should be called.

There are two ways that you might implement the calling of a C function (e.g. `add()`, `lw()`, `jal()`) based on the value of the opcode and, in the case of R-instructions, the `funct` field of an instruction:

- Use a switch statement for the opcode field, e.g.,

```
#define ADDI_OPCODE (0x08)
...
switch (opcode) {
    case 0:  handle_r_instructions(instruction);
            break;
    case ADDI_OPCODE: addi(instruction);
            break;
    ...
}
```

where `handle_r_instructions()` has a similar switch statement for calling the R-instruction function (`add()`, `sub()`, etc.) based on the `funct` field, or

- Use a dispatch table, which is an array of function pointers (see the lecture notes on function pointers from earlier in the semester). For example, you could declare the table `opcode_dispatch_table` as follows:

```
typedef void (*OP_DISPATCH) (uint64_t);

OP_DISPATCH opcode_dispatch_table[MAX_NUMBER_OF_OPCODES];
```

You would need need to populate the dispatch table with the functions for each instruction, indexed by the opcode:

```
#define ADDI_OPCODE (0x08)
...
opcode_dispatch_table[0] = dispatch_r_instruction;
opcode_dispatch_table[ADDI_OPCODE] = &addi;
...
```

where `dispatch_r_instruction()` would similarly call an R-instruction function (e.g. `add()`) based on the `funct` field (a separate dispatch table could be used to call the appropriate R-instruction function based on the value of the `funct` field).

When an instruction is to be executed, your `cpu_execute()` function would simply use the opcode as an index into the dispatch table, and call the function found at that index in the table.

7. As discussed above, a program to be executed is loaded into memory by the `main()` function already defined in `test_programs.c`. If your `cpu` executable is getting a different output than the `ben_cpu` executable, you can look at the instructions that are being executed in the file `test_programs.txt` (or even `test_programs.c`) to figure out which instruction your code isn't executing correctly.

In looking at `test_programs.txt`, notice that the register names are `$0`, `$1`, ..., `$31` and constants are written without dollar signs – the exact reverse of Intel assembly code (this has nothing to do with the code you need to write, though).